

Assignment 1: Indexing

Nicolai Dahl Blicher-Petersen and Daniel Jonathan Smith
{s3441163, s3361789}@student.rmit.edu.au

RMIT

1 Index Construction

When parsing the data set of hundreds of documents we search for the HEADLINE and TEXT tags, which contain content for each document. During this process the Document ID is noted when encountering the DOCID tag, and all other tags are skipped.

Document ID Assignment

Internal Document IDs are assigned sequentially as they are encountered in the data set using a separate DocIdHandler class, which maintains a dynamically sized array of "raw" Document IDs from the data set (implemented using the standard Java ArrayList collection). The internal Document ID is an integer corresponding to the array index of the "raw" Document ID within the DocIdHandler class, to allow constant time lookups.

Parsing

Parsing content character by character is performed in accordance with the following set of rules:

1. When reaching mark-up tags such as `<p>`, skip them
2. When parsing a token starting with a letter, only allow letters, numbers, and hyphens as non-terminators of the token. When reaching a terminator for a token containing a hyphen, the word is analyzed and handled in the following way:
 - (a) If there is only one hyphen in the word and the first part is a common prefix such as co, pre, meta, or multi, the hyphen is removed and the two parts are concatenated into one [7][4].
 - (b) If there is only one hyphen in the word and the last term ends on 'ed' (e.g. case-based) the hyphen is kept [7].
 - (c) Otherwise all hyphens in the token are removed and a token with n hyphens becomes $n + 1$ tokens with no hyphens.
3. When parsing a token starting with a digit, only allow letters, digits, dots, and commas as non-terminators of the token. If commas or dots are at the end of a token they are removed.
4. If the token is less than three characters long, it is thrown away.

Note that this means that the space character will most often be the terminator of a token and that we do not handle acronyms in any special way. As described by Manning et al. [3, p. 24] handling hyphens is often done in accordance with some heuristics that keep or remove hyphens based on various attributes of a token, e.g. its length. Our parser is simplified and will produce unwanted results for sequences such as "San Francisco-Los Angeles". It will, however, manage to index words like "co-operative/cooperative" under the same term ID.

Inverted Index Construction in Memory

After each document has been parsed, the tokens that were gathered are passed to the indexer module as a term list $(t, f_{d,t})$, where t denotes the term and $f_{d,t}$ denotes the term frequency, in this context the within-document frequency of the term.

The in-memory representation of the inverted index is a `HashMap` that maps each term t to a document frequency f_t and a list of postings of the form $(d, f_{d,t})$, where d denotes the Document ID and $f_{d,t}$ again denotes the within-document frequency of the term. This, in turn, is represented as an `ArrayList`, sorted by Document ID.

Since our postings list insertion method inserts new postings in sort order, it guarantees a constant-time insertion of a document where the Document ID is greater than all existing Document IDs, and a linear-time insertion in other cases. We consistently provide a constant-time insertion in our parser implementation, as Document IDs are assigned sequentially and therefore insertion into the postings list is always done at the end.

Initially we tried passing a raw term list for each document to the indexer. Aggregating the terms within the parser provides a dramatic performance improvement when adding the terms to the inverted index, however, as each term requires only a single insertion into the postings list.

Inverted Index Representation on Disk

The inverted index implementation was inspired by Manning et al. [3]. We produce a lexicon, where each entry consists of the term t , the document frequency f_t , the byte offset of the associated postings list in our inverted list file and the byte size of the postings list entry. The byte size is used for retrieving compressed data, as discussed in section 5. These fields are separated by the pipe character and each term entry is separated by a newline.

The inverted list is saved on disk as a binary file consisting of variable byte encoded integer values (see section 5, to be retrieved via the byte offset and byte size stored in the lexicon. If a term t occurs in document d there will be a posting of two integers $(d, f_{d,t})$ denoting the Document ID followed by the within-document frequency of t .

2 Stoplist

The stopping as well as the indexing module have been implemented using the delegation design pattern. The stopper module's sole responsibility is to check if a word passed to it is in the stoplist provided when calling the indexing program [1].

2.1 Hash Table

To perform constant-time lookups in the stoplist we used a standard Java `HashSet` of strings, which is internally represented as a hash table [5]. We utilize the standard Java string hash function to distribute keys between the buckets, as it has been shown to have an even distribution of hash values when used with random string values [2]. Since the size and average character distribution of the stoplist is not known in advance it is necessary to use a hash function that can provide a reasonably even distribution of hash values over a random collection of strings.

When collating terms, the parser simply checks the stopper for the existence of each term in the stoplist. If the term exists in the stoplist it is discarded and the parser moves on to the next term.

3 Index Search

Parsing of Query Terms

Our index search implementation first reads in the search terms from the command line and runs them through the same parser used to construct the inverted index. This is to ensure the same tokenisation and normalisation processes are applied to the search terms. Without this step search terms provided may not match the representation of an initially identical indexed term.

Reading in the Lexicon

The lexicon of the inverted index is read into a standard Java HashMap collection, mapping a term t to term data (f_t, p_t, s_t) , where f_t is the document frequency, p_t is the postings list address in the inverted list file, and s_t is the postings list size in bytes as it is stored in the inverted list file. The postings list address is the byte offset of the postings list in the inverted list file.

Implementing the lexicon as a HashMap provides a constant time lookup for each term's frequency data, postings list address and postings list size.

Reading in the Postings List

The inverted list file is opened for random access using a Java SeekableByteChannel, which allows us to seek to arbitrary points in the file and retrieve multiple bytes at a time[6].

Once a search term is parsed, the postings list address, postings list byte size and document frequency are retrieved from the lexicon hash.

To retrieve the postings list we first seek to the byte offset stored in the postings list address for the term in the inverted list file. Each posting is made up of two variable byte encoded integers. Due to the variable byte encoding the number of bytes to read from disk cannot be deduced from the document frequency of the term. To improve efficiency the byte size of the postings list is stored in our lexicon, allowing us to avoid multiple disk reads by reading all bytes for the postings list into a buffer for parsing.

The buffer is first decoded to produce a list of integers. This list is parsed to construct a postings list consisting of postings of the form $(d, f_{d,t})$, denoting the Document ID followed by the within-document frequency of t . The postings list is passed back to the query program as a search result for further processing.

Retrieving the Raw Document ID

The raw Document ID mapping is read from the map file into the DocIdHandler class referred to in Section 1. The DocIdHandler stores the raw Document IDs in an array, where the raw Document ID resides in the array index of its corresponding internal Document ID. The DocIdHandler handles this representation internally, and simply converts an internal Document ID to a raw Document ID when requested.

Once the postings list is generated, this list is stepped over. Each raw Document ID is retrieved from the DocIdHandler and printed to the command line with the within-document frequency.

4 Index Size

When indexing the sample latimes collection, our final inverted index (including the lexicon, inverted list and document ID map files) with stopwords removed is approximately 176MB, where the original collection is approximately 476MB.

This is due to several factors:

1. The original latimes collection includes markup tags and extra metadata about each document, which we discard.
2. We remove stopwords from the collection.
3. Where each term in the collection requires disk space for an entire ASCII representation of each occurrence of a term, our inverted list simply requires space for the byte representation of two integers for each document in which a term occurs. Although in the worst case, in which a term occurs only once in a single document, this in fact occupies more space on disk, in the vast majority of cases there is a significant space saving.

5 Compression

Compression of the index is achieved using a variable-byte coding algorithm to encode document IDs and frequencies before storage in the inverted list file. This is implemented from the algorithm described by Manning et al. [3, p. 96-98]

Integers are encoding using a variable byte encoding scheme, where the most significant bit is a continuation bit and the remaining bits of the byte encode part of the integer. The continuation bit is set to 0 for all but the last byte of the encoded integer, which is set to 1 to signal the end of the encoded value.

Encoded variable byte codes are decoded by reading a series of bytes with continuation bit 0 followed by a byte with continuation bit 1. The remaining 7 bits of each of these bytes are concatenated and converted to an integral value.

Using variable byte encoding makes it much more efficient to store smaller values. Since postings lists are sorted by document ID we can exploit this by storing the gaps between document IDs rather than the document IDs themselves. Aside from cases in which the document is first in the postings list, where by necessity the gap value is equal to the document ID, the gap value will always be smaller than the document ID itself.

Compressing the index is undertaken post-processing, when the lexicon and inverted list are written to disk. We have elected to store the size in bytes of the postings list for each term in the lexicon to allow us to read the byte list in buffer in one go and minimise disk reads.

The postings list for a term is decoded into an array of integers when it is read in from the inverted list file for query processing.

The combination of these two compression methods brings the size of our final inverted index size down by two thirds from approx. 176MB to approx. 57MB.

A Contributions

A.1 Nicolai Dahl

1. Parser
2. Stopper
3. Search term parsing
4. Processing and display of search results
5. Index/Search front end

A.2 Daniel Smith

1. Index construction
2. Index representation on disk
3. Hash function research
4. Reading lexicon and inverted list from file
5. Parsing inverted list to produce search results
6. Compression

A.3 Equal Contribution

1. Program architecture
2. Report
3. Testing and bug fixing

Bibliography

- [1] Delegation pattern. http://www.princeton.edu/~achaney/tmve/wiki100k/docs/Delegation_pattern.html, 2013.
- [2] Neil Coffey. http://www.javamex.com/tutorials/collections/hash_function_technical_2.shtml, title = How the String hash function works, 2012.
- [3] Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. *Introduction to Information Retrieval*. Cambridge University Press, New York, 2008.
- [4] Richard Nordquist. Common prefixes in english. <http://grammar.about.com/od/words/a/comprefix07.htm>, 2013.
- [5] Oracle. Class `hashset<e>`. <http://docs.oracle.com/javase/6/docs/api/java/util/HashSet.html>, 2011.
- [6] Oracle. Interface `seekablebytechannel`. [http://openjdk.java.net/projects/nio/javadoc/java/nio/channels/SeekableByteChannel.html#position\(long\)](http://openjdk.java.net/projects/nio/javadoc/java/nio/channels/SeekableByteChannel.html#position(long)), 2013.
- [7] Craig Trim. The art of tokenization. <https://www.ibm.com/developerworks/community/blogs/nlp/entry/tokenization?lang=en>, 2013.